NONPROFIT AI RESOURCES · MANDY HATHAWAY CONSULTING

AI Risk and Continuity for Nonprofits

What happens when things go wrong, and how to be ready

Mandy Hathaway Consulting · mandyhathaway.com · AI Ethics & Technology Consulting

"What if it breaks?" is one of the most common questions nonprofits ask about AI tools. It is a good question, and it deserves a real answer rather than reassurance. AI tools do fail. They produce wrong information. Vendors go down, change their terms, raise their prices, or shut down entirely. Any technology your organization depends on for its work carries some level of continuity risk.

The goal is not to avoid that risk entirely. It is to understand it clearly enough to make deliberate choices about which tools to rely on, how deeply to integrate them, and what your fallback looks like when something does not work as expected.

The organizations that handle AI failures well are not the ones where nothing goes wrong. They are the ones that planned for it before it happened.

The types of failure that actually happen

AI risk in a nonprofit context comes in several distinct forms, and they require different responses.

1. **The tool produces wrong information.**
AI tools hallucinate. They generate plausible-sounding content that is factually incorrect. This happens most often with specific facts, citations, dates, names, and regulatory details. If an AI-drafted grant report cites a statistic that does not exist, and it goes out under your organization's name, that is a credibility problem. Review is not optional for anything consequential.

2. **The tool goes down.**
Cloud-based AI tools experience outages. If your staff has built workflows that depend on a particular tool, an outage can disrupt a workday. This is the same risk you take with any cloud-based software, and the mitigation is the same: do not build single-point-of-failure dependencies, and know what manual process exists as a backup.

3. **The vendor changes terms or pricing.**
AI is a rapidly evolving market. Vendors have raised prices, changed free-tier limits, discontinued products, and altered data handling policies with relatively short notice. If your organization has built significant workflows around a tool that changes its pricing structure, you may face a real budget problem or a disruptive migration.

4. **The tool is used outside its appropriate scope.**
Staff use AI tools for tasks they were not designed for, or in ways that exceed what the tool can reliably do. An AI tool summarizing meeting notes is low risk. An AI tool being asked to make clinical or legal recommendations is high risk. Scope creep in AI use is a category of failure that organizations create themselves.

5. **The vendor shuts down or is acquired.**
Startup AI companies are acquired, merged, or shut down. If your organization has data stored with a vendor that closes, you need to know how to export it, and you need to have done that before the closure is announced.

How to reduce dependency risk

The most effective risk mitigation is also the simplest: do not let any single AI tool become so embedded in your operations that losing it would be a crisis. This does not mean avoiding AI. It means being thoughtful about how central any one tool becomes.

> A USEFUL TEST
>
> For any AI tool your staff relies on regularly, ask: if this tool disappeared tomorrow with 48 hours notice, what would break, and how long would it take us to recover? If the answer involves significant program disruption or data loss, your dependency has gotten ahead of your risk planning.

Practical steps to reduce dependency risk:

1. **Keep your data portable.**
Do not store anything in an AI platform that does not also exist somewhere you control. Documents, case notes, reports, and other organizational records should live in your own systems, with AI tools used to work on them, not to store them.

2. **Know your export options before you need them.**
   For any tool where your organization's data accumulates, find the export function before there is an urgency to use it. Knowing how to get your data out takes five minutes; not knowing can take days at the worst possible moment.

3. **Avoid single-vendor lock-in for critical workflows.**
   If a workflow is genuinely critical, build it around outputs and processes your staff controls, not around features unique to one vendor's platform.

4. **Keep a manual fallback for anything AI-assisted.**
   If a task was done manually before AI, the knowledge of how to do it that way should not disappear. Document the manual process, even if it is rarely needed.

Managing AI errors in your work

The failure mode most likely to affect a nonprofit is not a dramatic breach or a vendor shutdown. It is a staff member sending out AI-generated content without adequate review, and that content being wrong.

Errors cluster in predictable places. AI tools are most likely to produce unreliable output when asked about specific statistics, regulations, policy details, or citations; when working outside their training data; when the task requires nuanced judgment about a specific person or situation; and when asked to do things that look like they are in scope but are actually at the edges of the tool's capability.

A PRACTICAL REVIEW STANDARD

For internal use only (staff notes, internal summaries, drafts for further editing): a light read for obvious errors is sufficient.

For anything leaving the organization (funder reports, donor communications, public statements, client-facing materials): a full human review against source materials is required before it goes out.

For anything with regulatory or legal implications: no AI-generated content should be used without review by someone with relevant expertise.

Building a culture of appropriate reliance

The organizations that use AI most effectively are not the ones that trust it least or the ones that trust it most. They are the ones that have developed a calibrated sense of where AI is reliable and where it is not, and have built that into how staff approach the tools.

That calibration does not come from a policy document. It comes from experience, conversation, and a culture where staff feel comfortable flagging when something does not look right rather than sending it anyway. Creating space for staff to share AI failures without blame is one of the most practical things a leader can do to reduce risk over time.

When something goes wrong with an AI tool — an error reaches a funder, a tool loses data, a staff member uses a tool in a way that created a privacy problem — treat it as information rather than incident. Document what happened, understand why, adjust the practice, and move forward. The organizations that build resilience do it by learning from small failures before they become large ones.